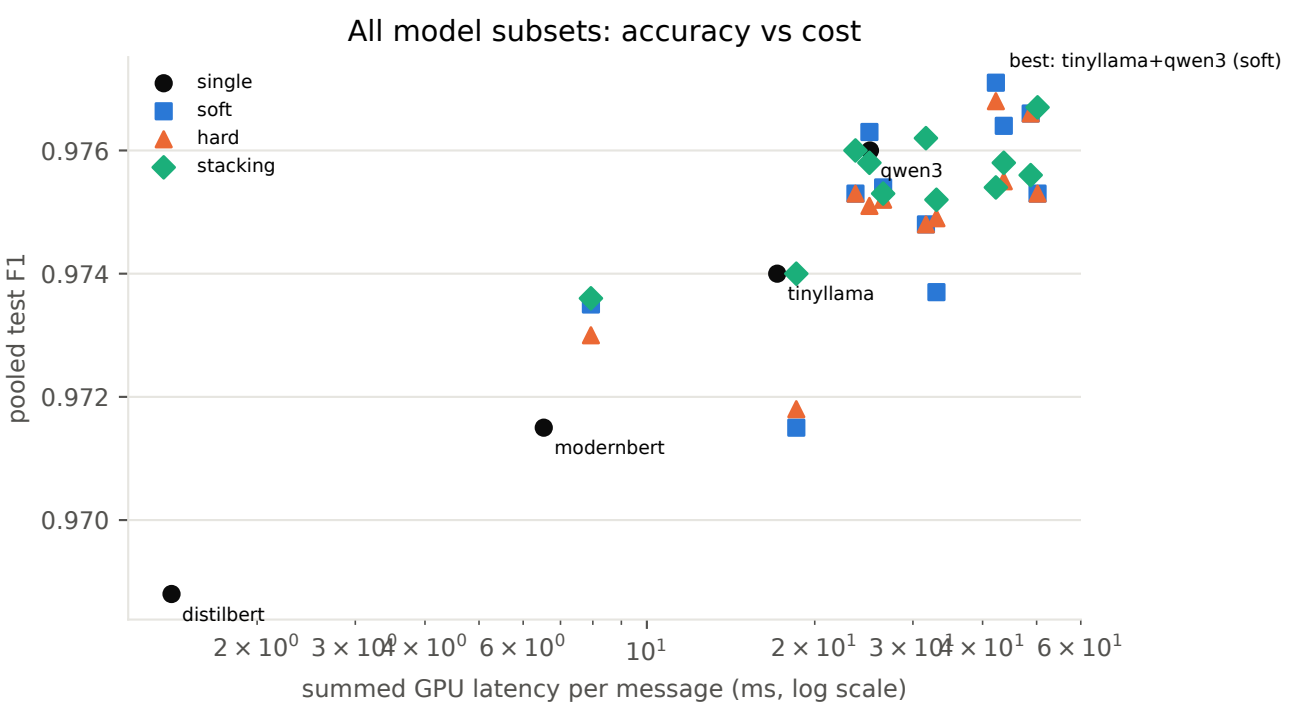
All model subsets: accuracy vs cost
single
soft
hard
stacking
best: tinyllama+qwen3 (soft)
qwen3
tinyllama
modernbert
distilbert
pooled test F1
0.976
0.974
0.972
0.970
summed GPU latency per message (ms, log scale)